

INT 334 Enterprise Application Automation

CA – 1

Name : Shubhansu Kumar Singh

Reg No : 12104991

Section : KO308

Roll : B59

Date : 19-02-2025

Practical : Create a sample Kubernetes deployment establishing a master -slave connection.

Create Shubhansu-Master EC2 instance

The screenshot displays the AWS Management Console for the 'Instances' page. The left sidebar shows navigation options like Dashboard, EC2 Global View, Events, and a list of instances. The main content area shows a table of instances with columns for Name, Instance ID, Instance state, Instance type, Status check, Alarm status, and Availability Zone. Two instances are listed: 'Shubhansu-Master' (i-05a14ce25e4ae79d4) and 'Singh-Slave' (i-0d357a0f066ef55a6), both in a 'Running' state. Below the table, the details for the 'Shubhansu-Master' instance are expanded, showing fields like Required, Operator, Instance details (AMI ID, AMI name), Stop protection, Instance auto-recovery, Monitoring, Allowed image, Launch time, Lifecycle, Platform details, Termination protection, AMI location, and Stop-hibernate behavior.

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone
Shubhansu-Master	i-05a14ce25e4ae79d4	Running	t2.medium	Initializing	View alarms +	us-east-1a
Singh-Slave	i-0d357a0f066ef55a6	Running	t2.micro	Initializing	View alarms +	us-east-1d

i-05a14ce25e4ae79d4 (Shubhansu-Master)

- Required: [arn:aws:ec2:us-east-1:600790316142:instance/i-05a14ce25e4ae79d4](#)
- Operator: -
- Instance details:
 - AMI ID: [ami-04b4f1a9cf54c11d0](#)
 - AMI name: [ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-amd64-server-20250115](#)
- Stop protection: Disabled
- Instance auto-recovery: Default
- Monitoring: disabled
- Allowed image: -
- Launch time: [Wed Feb 19 2025 12:25:11 GMT+0530 \(India Standard Time\)](#) (4 minutes)
- Lifecycle: normal
- Platform details: [Linux/UNIX](#)
- Termination protection: Disabled
- AMI location: [amazon/ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-amd64-server-20250115](#)
- Stop-hibernate behavior: Disabled

Create Singh-Slave EC2 instance

The screenshot displays the AWS Management Console for the 'Instances' page, showing the details for the 'Singh-Slave' instance. The left sidebar is the same as the previous screenshot. The main content area shows the instance details for 'i-0d357a0f066ef55a6 (Singh-Slave)'. The 'Details' tab is selected, showing fields like Instance ID, IP address, Hostname type, Answer private resource DNS name, Auto-assigned IP address, Public IPv4 address, Instance state, Private IP DNS name (IPv4 only), Instance type, VPC ID, Private IPv4 addresses, Public IPv4 DNS, and Elastic IP addresses.

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone
Shubhansu-Master	i-05a14ce25e4ae79d4	Running	t2.medium	Initializing	View alarms +	us-east-1a
Singh-Slave	i-0d357a0f066ef55a6	Running	t2.micro	Initializing	View alarms +	us-east-1d

i-0d357a0f066ef55a6 (Singh-Slave)

- Instance ID: [i-0d357a0f066ef55a6](#)
- IP address: -
- Hostname type: IP name: [ip-172-31-88-88.ec2.internal](#)
- Answer private resource DNS name: IPv4 (A)
- Auto-assigned IP address: [52.91.114.119](#) [Public IP]
- Public IPv4 address: [52.91.114.119](#) | [open address](#)
- Instance state: [Running](#)
- Private IP DNS name (IPv4 only): [ip-172-31-88-88.ec2.internal](#)
- Instance type: t2.micro
- VPC ID: [vpc-05f60cda716daf184](#)
- Private IPv4 addresses: [172.31.88.88](#)
- Public IPv4 DNS: [ec2-52-91-114-119.compute-1.amazonaws.com](#) | [open address](#)
- Elastic IP addresses: -
- AWS Compute Optimizer finding: [Opt-in to AWS Compute Optimizer for recommendations.](#)

Initialize Kubernetes cluster in master node

```
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgpfilters.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-cni-plugin created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-cni-plugin created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
[INFO] Master node setup complete!
[INFO] To join worker nodes, use the following command:
[INFO] Run 'cat join-command.sh' to see the join command.
root@ip-172-31-25-112:/home/ubuntu# cat join-command.sh
kubeadm join 172.31.25.112:6443 --token f0kf3q.suflofib0whi6x8c --discovery-token-ca-cert-hash sha256:602b97e2496dca7b50c77659f2dfce2d3e542419c94be3b4f235ebd476f7dd7
root@ip-172-31-25-112:/home/ubuntu#
```

i-05a14ce25e4ae79d4 (Shubhansu-Master)

PublicIPs: 54.162.125.6 PrivateIPs: 172.31.25.112

Join slave node to the cluster :

```
[reset] Stopping the kubelet service
[reset] Unmounting mounted directories in "/var/lib/kubelet"
[reset] Deleting contents of directories: [/etc/kubernetes/manifests /var/lib/kubelet /etc/kubernetes/pki]
[reset] Deleting files: [/etc/kubernetes/admin.conf /etc/kubernetes/super-admin.conf /etc/kubernetes/kubelet.conf /etc/kubernetes/bootstrap-kubelet.conf /etc/kubernetes/controller-manager.conf /etc/kubernetes/scheduler.conf]

The reset process does not clean CNI configuration. To do so, you must remove /etc/cni/net.d

The reset process does not reset or clean up iptables rules or IPVS tables.
If you wish to reset iptables, you must do so manually by using the "iptables" command.

If your cluster was setup to utilize IPVS, run ipvsadm --clear (or similar)
to reset your system's IPVS tables.

The reset process does not clean your kubeconfig files and you must remove them manually.
Please, check the contents of the $HOME/.kube/config file.
[INFO] Joining the Kubernetes cluster...
root@ip-172-31-88-88:/home/ubuntu# sudo kubeadm join 172.31.25.112:6443 --token f0kf3q.suflofib0whi6x8c --discovery-token-ca-cert-hash sha256:602b97e2496dca7b50c77659f2dfce2d3e542419c94be3b4f235ebd47a6f7dd7
[preflight] Running pre-flight checks/a6f7dd7
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
root@ip-172-31-88-88:/home/ubuntu#
```

i-0d357a0f066ef55a6 (Singh-Slave)

PublicIPs: 52.91.114.119 PrivateIPs: 172.31.88.88

Check node status in master node : Cluster ready

```
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-cni-plugin created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-cni-plugin created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
[INFO] Master node setup complete!
[INFO] To join worker nodes, use the following command:
[INFO] Run 'cat join-command.sh' to see the join command.
root@ip-172-31-25-112:/home/ubuntu# cat join-command.sh
kubeadm join 172.31.25.112:6443 --token f0kf3q.suflofib0whi6x8c --discovery-token-ca-cert-hash sha256:602b97e2496dca7b50c77659f2dfce2d3e542419c94be3b4f235ebd47a6f7dd7
root@ip-172-31-25-112:/home/ubuntu# kubectl get nodes
NAME        STATUS    ROLES    AGE     VERSION
master      Ready     control-plane  2m28s   v1.29.14
worker      NotReady  <none>      19s     v1.29.14
root@ip-172-31-25-112:/home/ubuntu# kubectl get nodes
NAME        STATUS    ROLES    AGE     VERSION
master      Ready     control-plane  2m40s   v1.29.14
worker      Ready     <none>      31s     v1.29.14
root@ip-172-31-25-112:/home/ubuntu#
```

i-05a14ce25e4ae79d4 (Shubhansu-Master)

PublicIPs: 54.162.125.6 PrivateIPs: 172.31.25.112


Create a sample deployment and check status

```
root@ip-172-31-25-112:/home/ubuntu# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-86dcdf4c6-66jhf    1/1     Running   0           2m5s
nginx-deployment-86dcdf4c6-869gb    1/1     Running   0           2m5s
nginx-deployment-86dcdf4c6-zjnfq    1/1     Running   0           2m5s
root@ip-172-31-25-112:/home/ubuntu# kubectl get nodes
NAME        STATUS    ROLES    AGE     VERSION
master      Ready     control-plane  8m6s    v1.29.14
worker      Ready     <none>      5m57s   v1.29.14
root@ip-172-31-25-112:/home/ubuntu# kubectl get services
NAME      TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes ClusterIP   10.96.0.1      <none>          443/TCP          8m10s
nginx     NodePort    10.110.174.33 <none>          80:30295/TCP     78s
root@ip-172-31-25-112:/home/ubuntu# kubectl cluster-info
Kubernetes control plane is running at https://172.31.25.112:6443
CoreDNS is running at https://172.31.25.112:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
root@ip-172-31-25-112:/home/ubuntu# kubectl get pods --all-namespaces
NAMESPACE   NAME                                READY   STATUS    RESTARTS   AGE
default     nginx-deployment-86dcdf4c6-66jhf    1/1     Running   0           2m54s
default     nginx-deployment-86dcdf4c6-869gb    1/1     Running   0           2m54s
default     nginx-deployment-86dcdf4c6-zjnfq    1/1     Running   0           2m54s
kube-system calico-kube-controllers-74d5f9d7bb-rjq65 1/1     Running   0           8m41s
kube-system calico-node-6rd2x              1/1     Running   0           6m40s
kube-system calico-node-bfgks            1/1     Running   0           8m41s
kube-system coredns-76f75df574-2wgs2    1/1     Running   0           8m40s
kube-system coredns-76f75df574-6dggp    1/1     Running   0           8m41s
kube-system etcd-master                 1/1     Running   0           8m45s
kube-system kube-apiserver-master        1/1     Running   0           8m45s
kube-system kube-controller-manager-master 1/1     Running   0           8m48s
kube-system kube-proxy-s6vpk             1/1     Running   0           8m41s
kube-system kube-proxy-xtvtl            1/1     Running   0           6m40s
kube-system kube-scheduler-master        1/1     Running   0           8m46s
```

i-05a14ce25e4ae79d4 (Shubhansu-Master)

PublicIPs: 54.162.125.6 PrivateIPs: 172.31.25.112



Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

The above mentioned screen shot shows the successful deployment of sample service in Kubernetes service.

The following commands were executed for this practical.

Setup MASTER-NODE Connection

```
#!/bin/bash
```

```
# Function to log messages
```

```
log() {  
    echo "[INFO] $1"  
}
```

```
# Set hostname (Modify as needed)
```

```
NODE_TYPE=$1
```

```
if [ "$NODE_TYPE" == "master" ]; then
```

```
    sudo hostnamectl set-hostname master
```

```
elif [ "$NODE_TYPE" == "worker" ]; then
```

```
    sudo hostnamectl set-hostname worker
```

```
else
```

```
    echo "Usage: $0 [master|worker]"
```

```
    exit 1
```

```
fi
```

```
# Update system
```

```
log "Updating system..."
```

```
sudo apt-get update && sudo apt-get upgrade -y
```

```
# Disable swap
```

```
log "Disabling swap..."
```

```
sudo swapoff -a
```

```
# Load necessary kernel modules
```

```
log "Loading required kernel modules..."
```

```
cat << EOF | sudo tee /etc/modules-load.d/k8s.conf
```

```
overlay
```

```
br_netfilter
```

```
EOF
```

```
sudo modprobe overlay
```

```
sudo modprobe br_netfilter
```

```
# Set sysctl parameters
```

```
log "Configuring networking..."
```

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOF
```

```
sudo sysctl --system
```

```
lsmod | grep br_netfilter
```

```
lsmod | grep overlay
```

```
# Install container runtime (containerd)
```

```
log "Installing container runtime..."
```

```
sudo apt-get update
```

```
sudo apt-get install -y ca-certificates curl
```

```
# Add Docker GPG key and repository
```

```
log "Adding Docker repository..."
```

```
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o  
/etc/apt/keyrings/docker.asc
```

```
sudo chmod a+r /etc/apt/keyrings/docker.asc
```

```
echo "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc]  
https://download.docker.com/linux/ubuntu $(. /etc/os-release && echo  
\"$VERSION_CODENAME\") stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

```
sudo apt-get update
```

```
sudo apt-get install -y containerd.io
```

```
# Configure containerd
```

```
log "Configuring containerd..."
```

```
containerd config default | sed -e 's/SystemdCgroup = false/SystemdCgroup = true/' -e  
's/sandbox_image = "registry.k8s.io/pause:3.6"/sandbox_image =  
"registry.k8s.io/pause:3.9"/' | sudo tee /etc/containerd/config.toml
```

```
sudo systemctl restart containerd
```

```
sudo systemctl status containerd --no-pager
```



```
# Install Kubernetes packages
```

```
log "Installing Kubernetes components..."
```

```
sudo apt-get update
```

```
sudo apt-get install -y apt-transport-https ca-certificates curl gpg
```

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.29/deb/Release.key | sudo gpg --dearmor -o  
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
```

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]  
https://pkgs.k8s.io/core:/stable:/v1.29/deb/ /' | sudo tee  
/etc/apt/sources.list.d/kubernetes.list
```

```
sudo apt-get update
```

```
sudo apt-get install -y kubelet kubeadm kubectl
```

```
sudo apt-mark hold kubelet kubeadm kubectl
```

```
# Initialize the Kubernetes cluster (Only on master node)
```

```
if [ "$NODE_TYPE" == "master" ]; then
```

```
    log "Initializing Kubernetes master node..."
```

```
    sudo kubeadm init
```

```
    mkdir -p "$HOME/.kube"
```

```
    sudo cp -i /etc/kubernetes/admin.conf "$HOME/.kube/config"
```

```
    sudo chown "$(id -u)": "$(id -g)" "$HOME/.kube/config"
```

```
# Apply Calico network plugin
```

```
log "Applying Calico networking..."
```

```
kubectl apply -f
https://raw.githubusercontent.com/projectcalico/calico/v3.26.0/manifests/calico.yaml
```

```
log "Master node setup complete!"

log "To join worker nodes, use the following command:"

kubeadm token create --print-join-command > join-command.sh

chmod +x join-command.sh

log "Run 'cat join-command.sh' to see the join command."

elif [ "$NODE_TYPE" == "worker" ]; then

    log "Resetting Kubernetes on worker node..."

    sudo kubeadm reset --force

    log "Joining the Kubernetes cluster..."

    # Run the join command (manual step)

    log "Run the join command from the master node."

fi
```

```
# For Master
```

```
# chmod +x setup_k8s.sh
# ./setup_k8s.sh master
# cat join-command.sh
```

```
# For Worker
```

```
# chmod +x setup_k8s.sh
```

```
# ./setup_k8s.sh worker
```

```
# Copy the join command from the master node and run it manually on the worker node.
```

Create Deployment

```
#!/bin/bash
```

```
# Function to log messages with spacing
```

```
log() {
```

```
    echo -e "\n[INFO] $1\n"
```

```
}
```

```
# Cleanup option - Placed at the beginning to avoid unnecessary execution
```

```
if [ "$1" == "cleanup" ]; then
```

```
    log "Deleting the Nginx deployment..."
```

```
    kubectl delete deployment nginx-deployment
```

```
    log "Deleting the Nginx service..."
```

```
    kubectl delete service nginx
```

```
    log "Cleanup complete!"
```

```
    exit 0
```

```
fi
```

```
# Define YAML file name and URL
```

```
YAML_FILE="nginx-deployment.yaml"
```

```
YAML_URL="https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/controllers/nginx-deployment.yaml"
```

```
# Download YAML file
```

```
log "Downloading Nginx deployment YAML..."
```

```
curl -s -o "$YAML_FILE" "$YAML_URL"
```

```
if [ ! -f "$YAML_FILE" ]; then
```

```
    echo -e "\n[ERROR] Failed to download $YAML_FILE. Exiting.\n"
```

```
    exit 1
```

```
fi
```

```
# Apply the YAML configuration
```

```
log "Applying the Nginx deployment..."
```

```
kubectl apply -f "$YAML_FILE"
```

```
# Wait for deployment to be ready
```

```
log "Waiting for the deployment to be ready..."
```

```
kubectl rollout status deployment/nginx-deployment
```

```
# Create a NodePort service for Nginx
```

```
log "Creating a NodePort service for Nginx..."
```

```
kubectl create service nodeport nginx --tcp=80:80
```

```
# Display service information
```

```
log "Fetching details of the created Nginx service..."
```

```
kubectl get svc nginx
```

```
# Display deployment and cluster info
```

```
log "Listing all running pods..."
```

```
kubectl get pods
```

sleep 1

log "Listing all deployments..."

kubectl get deployment

sleep 1

log "Listing all services..."

kubectl get service

sleep 1

log "Fetching cluster information..."

kubectl cluster-info

sleep 1

log "Listing all pods across all namespaces..."

kubectl get pods --all-namespaces

log "Nginx deployment setup complete!"